

Request
with SLOs



TTFT SLO:100ms
TPOT SLO:150ms



TTFT SLO:3000ms
TPOT SLO:30ms




TTFT SLO:800ms
TPOT SLO:100ms

Temporal Deadline-Aware Request Dispatcher

Earliest-deadline-first admission

earliest




Waiting
Queue

③

Deadline-aware preemption

earliest




Running
set

④

Spatial Step-Time-Model-Guided Token Allocator

⑤ Active Decode Workload




KV Cache

⑥ Step Time Model




⑦ Safe token allocate



⑧

⑨

Higher muti-SLO
Goodput



LLM

Continuous batch and LLM Engine Step